

# Parallel Programming

# Different levels of parallelism

- Parallelism manifests itself at different **granularity** levels
  - **bit-level** parallelism
    - processing **multiple bits** of data in parallel
  - **instruction-level** parallelism
    - executing **different instructions** from the same instruction stream in parallel
  - **task-level** parallelism
    - executing **separate instruction streams** in parallel
- We will focus on programming techniques for **task-level parallelism**

# Classes of Parallel Computers

- Many different forms of parallel hardware
  - multi-core processors
  - symmetric multiprocessors
  - general purpose graphics processing unit
  - field-programmable gate arrays
  - computer clusters
- We will focus (in this section) on programming techniques for multi-cores and symmetric multiprocessors
  - GPUs and FPGAs have their own programming frameworks (eg. CUDA, OpenCL) and languages (eg. HDLs like Verilog)
- We will consider cluster computing in the last part of our course

# JVM parallelism

- JVM is based on thread-based parallelism
  - Each process can contain multiple independent concurrency units called threads
  - Threads can be started from within the same program, and they share the same memory address space
  - Each thread has a program counter and a program stack
  - JVM threads cannot modify each other's stack memory

# Creating, starting and waiting for threads

- Each JVM process starts with a **main thread**
- To start additional threads:
  - 1. Define a **Thread subclass**
  - 2. Instantiate a **new** Thread object
  - 3. Call **start** on the Thread object
- The Thread subclass defines, in a **run** method, the **code** that the thread will execute

```
class HelloThread extends Thread {  
    override def run() = {  
        println("Hello world!")  
    }  
}  
  
val t1 = new HelloThread; val t2 = new HelloThread  
t1.start(); t2.start()  
t1.join(); t2.join()
```

# Memory model

- Memory model is a set of rules that describes how threads interact when accessing shared memory
- The memory model for the JVM:
  - Two threads writing to separate locations in memory do not need synchronization
  - In order for a thread *X* to observe all the writes by thread *Y*, it is necessary to add synchronization points (e.g. *X* calls *join* on *Y* before reading)

```
val vector = Array.fill(10000)(0)
class th extends Thread {
  override def run() = {
    for (i <- 0 until 10000) {vector(i)=i; Thread.sleep(1)}
  }
}
val thread = new th; thread.start; Thread.sleep(5000)
vector.max
```

# Parallel computation

- Assume to have a **parallel** function (implemented using concurrency libraries out of the scope of this course):  
**def** parallel[A, B](taskA: => A, taskB: => B): (A, B)
  - that evaluates two expressions in parallel, and returns the **pair** composed of the two results
- Assume you have to compute the **p-norm** of a vector:

$$||a||_p := \left( \sum_{i=0}^{N-1} [|a_i|^p] \right)^{1/p}$$

- It is possible to compute the summation in **parallel** on distinct **segments** of the vector

# Parallel computation of the norm

```
def sumSegment(a: Array[Double], p: Double, s: Int, t: Int):  
Double = {  
  var i = s; var sum: Double = 0  
  while (i < t) {  
    sum = sum + Math.pow(a(i), p)  
    i = i + 1  
  }  
  sum  
}  
  
def pNormTwoParts(a: Array[Double], p: Double): Double = {  
  val m = a.length / 2  
  val (sum1, sum2) = parallel(sumSegment(a, a.length, 0, m),  
    sumSegment(a, p, m, a.length))  
  Math.pow((sum1 + sum2), 1 / p)  
}
```



# Unbounded parallel computation

```
def segmentRec(a: Array[Double], p: Double, s: Int, t: Int):  
Double = {  
  if (t - s < threshold)  
    sumSegment(a, p, s, t) // small segment done sequentially  
  else {  
    val m = s + (t - s) / 2  
    val (sum1, sum2) = parallel(  
      segmentRec(a, p, s, m),  
      segmentRec(a, p, m, t))  
    sum1 + sum2  
  }  
}  
  
def pNormRec(a: Array[Double]): Double =  
  Math.pow (segmentRec(a, a.length, 0, a.length), 1 / a.length)
```

# Parallel computation: what happens?


- Efficient parallelism requires support from
  - hardware
  - operating system
  - virtual machine (as in the case of JVM)
  - language
- Given sufficient resources, computation is parallelised and the parallel program can run faster

# Hardware architecture matters

- Assume we consider the same pattern of computation, **but only summing** without applying power
- The speed-up does not always increase, why?
  - The computation requires more frequent access to the **shared memory**
  - Memory, in this case, becomes a **bottleneck**

# Example: estimation of pi

- Consider a square of edge one and a circle of radius one with center in the bottom-left corner of the square:



- The ratio between the surface of  $\frac{1}{4}$  of the circle and the square:

$$\pi / 4$$

- Estimation of pi:
  - random sample points inside the square
  - count the frequency of those that fall inside the circle
  - multiply it by 4

# Sequential implementation

```
import scala.util.Random

def mcCount(iter: Int): Int = {
  val randomX = new Random
  val randomY = new Random
  var hits = 0
  for (i <- 0 until iter) {
    val x = randomX.nextDouble // in [0,1]
    val y = randomY.nextDouble // in [0,1]
    if (x * x + y * y < 1) hits = hits + 1
  }
  hits
}

def monteCarloPiSeq(iter: Int): Double =
  4.0 * mcCount(iter) / iter
```

# Parallel implementation

- We can split the computation in **parallel tasks**:
  - Here is a parallel implementation with **four** parallel tasks

```
def monteCarloPiPar(iter: Int): Double = {  
  val (pi1, pi2, pi3, pi4) = parallel(  
    mcCount(iter / 4),  
    mcCount(iter / 4),  
    mcCount(iter / 4),  
    mcCount(iter / 4)  
  )  
  4.0 * (pi1 + pi2 + pi3 + pi4) / iter  
}
```

# Exercise

- Write a parallel implementation of MergeSort

# Exercise

- Write a parallel implementation of MergeSort

```
def parMergeSort(xs: Array[Int], maxDepth: Int): Unit = {  
  val ys = new Array[Int](xs.length)  
  
  def sort(from: Int, until: Int, depth: Int): Unit = {  
    if (depth == maxDepth) {  
      quickS(xs, from, until)  
    } else {  
      val mid = (from + until) / 2  
      parallel(sort(mid, until, depth + 1),  
              sort(from, mid, depth + 1))  
      val flip = (maxDepth - depth) % 2 == 0  
      val src = if (flip) ys else xs  
      val dst = if (flip) xs else ys  
      merge(src, dst, from, mid, until)  
    }  
  }  
  ...  
}
```



# How fast are parallel programs?

- Parallel versions of algorithms can be faster:
  - Can we perform an analytical analysis of the gain, following an approach like asymptotic complexity for sequential algorithms?
- Time complexity  $T(n)$  of sequential algorithms quantifies the expected number of elementary operations to execute for an input of size  $n$ 
  - Such operations are executed in sequence (one at a time)
- For parallel algorithms, we use two measures:
  - Work complexity:  $W(n)$  (the same as the above  $T(n)$ )  
Expected number of elementary operations for an input of size  $n$
  - Depth complexity:  $D(n)$   
Expected length of the longest sequence of elementary operations that must be executed in sequence for an input of size  $n$

# Work and depth complexity

- We consider the `sumSegment` example seen before:

```
def sumSegment(a: Array[Double], p: Double, s: Int, t: Int):  
Double = {  
  var i = s; var sum: Double = 0  
  while (i < t) {  
    sum = sum + Math.pow(a(i), p)  
    i = i + 1  
  }  
  sum  
}
```

Time complexity in the sequential case:  $T(t-s) = c_1(t-s) + c_2$


# Work and depth complexity

- **Work complexity** of the parallel recursive version:

```
def segmentRec(a: Array[Double], p: Double, s: Int, t: Int):  
Double = {  
  if (t - s < threshold)  
    sumSegment(a, p, s, t)  
  else {  
    val m = s + (t - s) / 2  
    val (sum1, sum2) = parallel(  
      segmentRec(a, p, s, m),  
      segmentRec(a, p, m, t))  
    sum1 + sum2  
  }  
}
```

$$W(s, t) = \begin{cases} c_1(t - s) + c_2, & \text{if } t - s < \text{threshold} \\ W(s, m) + W(m, t) + c_3 & \text{otherwise, for } m = \lfloor (s + t)/2 \rfloor \end{cases}$$


# Work and depth complexity



- Assume  $s, t$  for which exists  $N$  s.t.  $t - s = 2^N(\text{threshold} - 1)$ 
  - $N$  depth of the tree, having  $2^N$  leaves and  $2^N - 1$  internal nodes
  - $W(s, t) = 2^N(c_1(\text{threshold} - 1) + c_2) + (2^N - 1)c_3 = 2^N c_4 + c_5$

$$W(s, t) = \begin{cases} c_1(t - s) + c_2, & \text{if } t - s < \text{threshold} \\ W(s, m) + W(m, t) + c_3 & \text{otherwise, for } m = \lfloor (s + t)/2 \rfloor \end{cases}$$

# Work and depth complexity



- Given any  $s, t$ , take  $N$  s.t.

$$2^{N-1}(\text{threshold}-1) < (t-s) \leq 2^N(\text{threshold}-1)$$

- As complexity increases, taking  $s', t'$  s.t.  $t'-s' = 2^N(\text{threshold}-1)$  we have  $W(s, t) \leq W(s', t') = 2^N c_4 + c_5$
- From  $2^{N-1}(\text{threshold}-1) < (t-s)$  we have  $2^N < 2(t-s) / (\text{threshold}-1)$ , hence we have  $W(s, t) < 2(t-s) / (\text{threshold}-1) c_4 + c_5$
- This means that  $W(s, t)$  is in  $O(t-s)$

# Work and depth complexity

- Depth complexity of the parallel recursive version:

```
def segmentRec(a: Array[Double], p: Double, s: Int, t: Int):  
Double = {  
  if (t - s < threshold)  
    sumSegment(a, p, s, t)  
  else {  
    val m = s + (t - s) / 2  
    val (sum1, sum2) = parallel(  
      segmentRec(a, p, s, m),  
      segmentRec(a, p, m, t))  
    sum1 + sum2  
  }  
}
```

$$D(s, t) = \begin{cases} c_1(t - s) + c_2, & \text{if } t - s < \text{threshold} \\ \max(D(s, m), D(m, t)) + c_3 & \text{otherwise, for } m = \lfloor (s + t)/2 \rfloor \end{cases}$$

# Work and depth complexity

- Assume  $s, t$  for which exists  $N$  s.t.  $t-s = 2^N(\text{threshold}-1)$ 
  - $N$  depth of the tree, having  $2^N$  leaves and  $2^N-1$  internal nodes
  - The value of  $D(s,t)$  in leaves is  $c_1(\text{threshold}-1) + c_2$
  - One level above is:  $c_1(\text{threshold}-1) + c_2 + c_3$
  - In the root,  $N$  levels above:  $c_1(\text{threshold}-1) + c_2 + N c_3 = N c_3 + c_4$
- Assume  $N$  s.t.  $2^{N-1}(\text{threshold}-1) < (t-s) \leq 2^N(\text{threshold}-1)$ 
  - As complexity increases, we have  $D(s,t) \leq N c_3 + c_4$
  - From  $2^{N-1}(\text{threshold}-1) < (t-s)$  we have  
 $N-1 < \log((t-s)/(\text{threshold}-1))$ , hence  $N < \log((t-s) / c_5) + 1$
  - Finally  $D(s,t) < (\log((t-s) / c_5) + 1) c_3 + c_4$  which is in  $O(\log(t-s))$

$$D(s, t) = \begin{cases} c_1(t-s) + c_2, & \text{if } t-s < \text{threshold} \\ \max(D(s, m), D(m, t)) + c_3 & \text{otherwise, for } m = \lfloor (s+t)/2 \rfloor \end{cases}$$

# Number of physically parallel threads

- Let  $P$  be the number of maximal parallel threads
- Can we bound running time as a function of  $P$ ?

$$D(e) + \frac{W(e)}{P}$$

- given the amount of work  $W(e)$ , at least  $W(e)/P$  time is needed
  - if  $P$  goes to infinity, the depth  $D(e)$  any way remains
- Observations:
  - if  $P$  is fixed the complexity grows as the complexity of the sequential solution (it is fine to assume  $D(e)$  bound by  $W(e)$ )
  - the existence of limits to the gain of parallelization (due to the presence of  $D(e)$ ) is known as Amdahl's Law



# Amdahl's law

- Suppose a task is divided in two parts:
  - part1: a fraction  $f$  that cannot be speed-up
  - part2: the remaining  $1-f$  that can be speed-up
- Speed-up: ratio between the sequential and the parallel execution time
- If we make part2  $P$  times faster the speed-up is:

$$1 / \left( f + \frac{1 - f}{P} \right)$$

- For example, if  $P=100$  and  $f=0.4$  we obtain 2.46
  - even if we speed-up the second part infinitely, we can obtain at most a global speed-up of  $1/0.4 = 2.5$

# Higher-order functions and parallelism

- Example: **sequential** and **parallel** version of map on Array

```
def mapArr[A,B](inp: Array[A], left: Int, right: Int,
                f : A => B, out: Array[B]): Unit = {
  var i= left
  while (i < right) {
    out(i)= f(inp(i))
    i= i+1
  }
}

def mapArrPar[A,B](inp: Array[A], left: Int, right: Int,
                  f : A => B, out: Array[B]): Unit = {
  if (right - left < threshold)
    mapArr(inp, left, right, f, out)
  else {
    val mid = left + (right - left)/2
    parallel(mapArrPar(inp, left, mid, f, out),
             mapArrPar(inp, mid, right, f, out))
  }
}
```

# Parallel map

- Parallelization on lists is **inconvenient**:
  - Split takes linear time (depth complexity is already linear!)
- Definition of a parallel **map** for the following trees
  - where the arrays in the leafs represent the partitions of the data in a collection
  - and the tree structure represents the order of split/combination of these partitions

```
trait Tree[A] { val size: Int }  
case class Leaf[A](a: Array[A]) extends Tree[A] {  
  override val size = a.size  
}  
case class Node[A](l: Tree[A], r: Tree[A]) extends Tree[A] {  
  override val size = l.size + r.size  
}
```

# Parallel map on trees

- Definition of a parallel **map** for the following trees

```
def mapTreePar[A,B:Manifest](t:Tree[A], f:A => B): Tree[B] =  
  t match {  
    case Leaf(a) => {  
      val len = a.length; val b = new Array[B](len)  
      var i = 0  
      while (i < len) { b(i)= f(a(i)); i= i + 1 }  
      Leaf(b) }  
    case Node(l,r) => {  
      val (lb,rb) = parallel(mapTreePar(l,f),  
                             mapTreePar(r,f))  
      Node(lb, rb) }  
  }
```

- Question: what is the **depth** complexity?

# Parallel map on trees

- Definition of a parallel **map** for the following trees

```
def mapTreePar[A,B:Manifest](t:Tree[A], f:A => B): Tree[B] =  
  t match {  
    case Leaf(a) => {  
      val len = a.length; val b = new Array[B](len)  
      var i = 0  
      while (i < len) { b(i)= f(a(i)); i= i + 1 }  
      Leaf(b) }  
    case Node(l,r) => {  
      val (lb,rb) = parallel(mapTreePar(l,f),  
                             mapTreePar(r,f))  
      Node(lb, rb) }  
  }
```

- **Question:** what is the **depth** complexity?
  - **Answer:** the **height** of the tree (hence logarithmic)

# Parallel Fold / Reduce

- We now move to **fold/reduce** higher order operations that apply a given operation to the elements in a collection:

```
List(1,3,8).foldLeft(100)((s,x) => s - x) ==  
  ((100 - 1) - 3) - 8 == 88
```

```
List(1,3,8).foldRight(100)((s,x) => s - x) ==  
  1 - (3 - (8 - 100)) == -94
```

```
List(1,3,8).reduceLeft((s,x) => s - x) ==  
  (1 - 3) - 8 == -10
```

```
List(1,3,8).reduceRight((s,x) => s - x) ==  
  1 - (3 - 8) == 6
```

# Towards a parallel reduce

- We consider a parallel implementation of **reduce** on the previous **trees**

```
def reduceTreePar[A](t: ArrayTree[A], f: (A,A) => A) : A =  
  t match {  
    case ArrayLeaf(a) => {  
      a.reduce(f)  
    }  
    case ArrayNode(l,r) => {  
      val (lr,rr) = parallel(reduceTreePar(l,f),  
                             reduceTreePar(r,f))  
      f(lr, rr) }  
  }
```

# Does the structure of the tree matters?

- Consider two trees with the same elements (in the same order) but with **different structure**:

```
def t1 =  
  ArrayNode[Int](  
    ArrayLeaf(Array(1)),  
    ArrayNode(ArrayLeaf(Array(3)), ArrayLeaf(Array(8)))  
  )  
  
def t2 =  
  ArrayNode[Int](  
    ArrayNode(ArrayLeaf(Array(1)), ArrayLeaf(Array(3))),  
    ArrayLeaf(Array(8))  
  )
```



# Does the structure of the tree matters?

- What are the values of the following expressions?

```
reduceTreePar(t1, (x:Int, y:Int) => x + y)  
reduceTreePar(t2, (x:Int, y:Int) => x + y)
```

```
reduceTreePar(t1, (x:Int, y:Int) => x - y)  
reduceTreePar(t2, (x:Int, y:Int) => x - y)
```

- Why **the structure matters** only in the first case?

# Does the structure of the tree matters?

- What are the values of the following expressions?

```
reduceTreePar(t1, (x:Int, y:Int) => x + y)  
reduceTreePar(t2, (x:Int, y:Int) => x + y)
```


```
reduceTreePar(t1, (x:Int, y:Int) => x - y)  
reduceTreePar(t2, (x:Int, y:Int) => x - y)
```

- Why **the structure matters** only in the first case?
  - Because minus is **not associative**, while plus is

# Associativity

- **Associativity:**

- an operation  $f: (A,A) \Rightarrow A$  is associative iff for every  $x, y, z$ :  
$$f(x, f(y, z)) = f(f(x, y), z)$$
- If we write  $f(a, b)$  in infix form as  $a \otimes \dots$ , associativity becomes  
$$x \otimes (y \otimes z) = (x \otimes y) \otimes z$$
- Graphically, using a **tree** representation:



# Property of associative operators


- Consider the following tree **visit**:

```
def toList[A](t: Tree[A]): List[A] = t match {  
  case Leaf(v) => List(v)  
  case Node(l, r) => toList(l) ++ toList(r) }
```

- Property of associativity:

given  $f(A,A) \Rightarrow A$  associative and  $t1:Tree[A]$  and  $t2:Tree[A]$  such that  $toList(t1) == toList(t2)$ , then  $reduce(t1,f) == reduce(t2,f)$

- Proof:** applying **rotation**, trees can be put in list-like normal form



- Note: rotation preserves **toList** and **reduce** (for associativity)

# Attention to associativity: floating point arithmetic

- Consider floating points values
  - Addition is **not associative!**

```
val e = 1e-200  
val x = 1e200  
val mx = -x
```

```
(x + mx) + e == x + (mx + e)
```

- Multiplication is **not associative!**

```
val y = 1e-200  
val z = 1e200
```

```
(y*z)*z == y*(z*z)
```

# Associative operations on tuples

- Suppose  $f_1: (A_1, A_1) \Rightarrow A_1$  and  $f_2: (A_2, A_2) \Rightarrow A_2$  are associative
  - Then  $f: ((A_1, A_2), (A_1, A_2)) \Rightarrow (A_1, A_2)$  defined by  $f((x_1, x_2), (y_1, y_2)) = (f_1(x_1, y_1), f_2(x_2, y_2))$  is also associative:
$$\begin{aligned} & f(f((x_1, x_2), (y_1, y_2)), (z_1, z_2)) = \\ & f((f_1(x_1, y_1), f_2(x_2, y_2)), (z_1, z_2)) = \\ & (f_1(f_1(x_1, y_1), z_1), f_2(f_2(x_2, y_2), z_2)) = (f_1, f_2 \text{ are associative}) \\ & (f_1(x_1, f_1(y_1, z_1)), f_2(x_2, f_2(y_2, z_2))) = \\ & f((x_1, x_2), (f_1(y_1, z_1), f_2(y_2, z_2))) = \\ & f((x_1, x_2), f((y_1, y_2), (z_1, z_2))) \end{aligned}$$
- We can also construct associative operations for  $n$ -tuples

# Example: average computation

- Given a collection of integers, compute the **average**

```
val sum = reduce(collection, _ + _)
val length = reduce(map(collection, (x:Int) => 1), _ + _)
sum/length
```

- This includes two reductions.  
Is there a solution using a **single reduce**?
- **Solution**: use pairs that compute **sum** and **length** at once

```
def f((sum1, len1), (sum2, len2)) = (sum1 + sum2, len1 + len2)
```

**f** is associative because addition is associative

```
val (sum, length) = reduce(map(collection, (x:Int)=>(x,1)), f)
sum/length
```

# Reduce on Arrays

- How to implement a **parallel reduce** higher-order function on **Arrays**?
  - convert the Array into a **balanced tree** (to reduce depth complexity)
  - do **tree reduction**
- Attention:
  - works only for **associative operations**, for which we can choose any tree representation that preserves the order of elements
- Optimization:
  - It is **not necessary** to actually construct the tree
  - It is sufficient to **apply directly the operator** instead of the Node constructor



# Reduce on Arrays

- How to implement a **parallel reduce** higher-order function on **Arrays**?

```
def reduceSeg[A](inp: Array[A], left: Int, right: Int,
                 f: (A,A) => A): A = {
  if (right - left < threshold) {
    var res = inp(left); var i = left+1
    while (i < right) { res = f(res, inp(i)); i = i+1 }
    res
  } else {
    val mid = left + (right - left)/2
    val (a1,a2) = parallel(reduceSeg(inp, left, mid, f),
                          reduceSeg(inp, mid, right, f))
    f(a1,a2)
  }
}

def reduce[A](inp: Array[A], f: (A,A) => A): A =
  reduceSeg(inp, 0, inp.length, f)
```

# Exercise

- Express the core of the **computation of norm** in terms of some fold/reduce higher-order function

$$\sum_{i=s}^{t-1} [|a_i|^p]$$

# Exercise

- Express the core of the **computation of norm** in terms of some fold/reduce higher-order function

$$\sum_{i=s}^{t-1} [|a_i|^p]$$

`foldleft(0)((s,x)=>s+pow(abs(x),p))`

- Question:** is it reasonable to expect a parallel implementation of `foldLeft`, like that of `reduce` ?

# Exercise

- Express the core of the **computation of norm** in terms of some fold/reduce higher-order function

$$\sum_{i=s}^{t-1} [|a_i|^p]$$

`foldleft(0)((s,x)=>s+pow(abs(x),p))`

- Question:** is it reasonable to expect a parallel implementation of `foldLeft`, like that of `reduce` ?
  - NO: `foldLeft` is intrinsically sequential, due to the use of the accumulator

# Aggregate

- For parallel higher-order programming, fold-like operations are replaced by alternative functions like **aggregate**

- On collections containing **A** objects, it is declared as:

```
def aggregate[B](z: =>B)(seqop: (B, A) => B, combop: (B, B) => B): B
```

- How to **implement** it?
    - See next slide

# Aggregate on Arrays

```
def aggregateSeg[A,B](inp: Array[A], left: Int, right: Int,
a:B, f: (B,A) => B, g: (B,B) => B): B = {
  if (right - left < threshold) {
    var res = a; var i = left
    while (i < right) { res = f(res, inp(i)); i = i+1 }
    res
  } else {
    val mid = left + (right - left)/2
    val (a1,a2) = parallel(
      aggregateSeg(inp, left, mid, a, f, g),
      aggregateSeg(inp, mid, right, a, f, g))
    g(a1,a2)
  }
}

def aggregatePar[A,B](inp: Array[A],
                      a:B, f: (B,A) => B, g: (B,B) => B): B =
  aggregateSeg(inp, 0, inp.length, a, f, g)
```

# Exercise

- Express the core of the **computation of norm** in terms of the **aggregate** higher-order function

$$\sum_{i=s}^{t-1} [|a_i|^p]$$

# Exercise

- Express the core of the **computation of norm** in terms of the **aggregate** higher-order function

$$\sum_{i=s}^{t-1} [|a_i|^p]$$

`aggregate(0) (_+pow(abs(_),p), _+_)`



# Towards a parallel scanLeft

- Higher-order functions include `scanLeft` and `scanRight`

```
val l = List(1,3,8)
l.scanLeft(100)(_ + _) //List(100, 101, 104, 112)
l.scanRight(100)(_ + _) //List(112, 111, 108, 100)
```

- We study the problem of implementing in `parallel` `scanLeft`
  - `scanRight` can be implemented `symmetrically`
  - let's start with a `sequential` implementation (on arrays)

```
def scanLeft[A](inp: Array[A], a0: A, f: (A,A) => A,
                out: Array[A]): Unit = {
  out(0) = a0; var a = a0; var i = 0
  while (i < inp.length) {
    a = f(a,inp(i)); i = i + 1; out(i)= a
  }
}
```

# Towards a parallel scanLeft


- In principle, we can use parallel **map** and **reduce** to obtain a parallel implementation with **logarithmic** depth

```
def mapSegS[A,B](inp: Array[A], left: Int, right: Int,
                  fi: (Int) => B, out: Array[B]): Unit = {
  if (right - left < threshold) {
    var i=left
    while (i < right) {
      out(i)=fi(i); i+=1
    }
  } else {
    val mid = left + (right - left)/2
    val (a1,a2) = parallel(mapSegS(inp, left, mid, fi, out),
                           mapSegS(inp, mid, right, fi, out))
  }
}

def scanLeftMR[A](inp: Array[A], a0: A, f: (A,A) => A, out: Array[A]) = {
  val fi = { (i:Int) => f(a0,reduceSeg(inp, 0, i, f)) }
  mapSegS(inp, 1, inp.length+1, fi, out)
  out(0) = a0
}
```

# Towards a parallel scanLeft


- Why is the **map-reduce** implementation not **satisfactory**?
  - Because the amount of **work** is **quadratic**!
  - Can we **avoid** such complexity? Yes, **storing** intermediary results



- With linear work store the reduce on array **segments** in a **tree**

# Towards a parallel scanLeft

- Given the tree in the previous slide compute the reduce for **each prefix** of the array as follows:



- along the path from the root to the segment, combine values in the **left-child** of nodes when move on the **right**
- **recompute** on the remaining values in the segment

# Towards a parallel scanLeft

- **First phase:** compute the tree with the intermediary results

```
trait TreeRes[A] { val res: A }
case class Leaf[A](from: Int, to: Int, override val res: A)
    extends TreeRes[A]
case class Node[A](l: TreeRes[A], override val res: A, r: TreeRes[A])
    extends TreeRes[A]

def upsweep[A](inp: Array[A], from: Int, to: Int,
    f: (A,A) => A): TreeRes[A] = {
  if (to - from < threshold)
    Leaf(from, to, reduceOnSeg(inp, from + 1, to, inp(from), f))
  else {
    val mid = from + (to - from)/2
    val (tL,tR) = parallel(upsweep(inp, from, mid, f),
        upsweep(inp, mid, to, f))
    Node(tL, f(tL.res,tR.res), tR)
  }
}
```

# Towards a parallel scanLeft

- **First phase:** compute the tree with the intermediary results

```
def reduceOnSeg[A](inp: Array[A], left: Int, right: Int,  
                  a0: A, f: (A,A) => A): A = {  
  var a = a0  
  var i = left  
  while (i < right) {  
    a = f(a, inp(i))  
    i = i+1  
  }  
  a  
}
```

# Towards a parallel scanLeft

- **Second phase:** go down the tree to compute scanLeft

```
def downsweep[A](inp: Array[A], a0: A, f: (A,A) => A, t: TreeRes[A],  
                 out: Array[A]): Unit = t match {  
  case Leaf(from, to, res) =>  
    scanLeftOnSeg(inp, from, to, a0, f, out)  
  case Node(l, _, r) => {  
    val (_,_) = parallel(downsweep(inp, a0, f, l, out),  
                        downsweep(inp, f(a0,l.res), f, r, out))  
  }  
}  
  
def scanLeftOnSeg[A](inp: Array[A], left: Int, right: Int, a0: A,  
                    f: (A,A) => A, out: Array[A]) = {  
  if (left < right) {  
    var i = left  
    var a = a0  
    while (i < right) {  
      a = f(a,inp(i))  
      i = i+1  
      out(i) = a  
    }  
  }  
}
```

# Towards a parallel scanLeft

- **Third phase:** put everything together

```
def scanLeftP[A](inp: Array[A],  
                 a0: A, f: (A,A) => A,  
                 out: Array[A]) = {  
  val t = upsweep(inp, 0, inp.length, f)  
  downsweep(inp, a0, f, t, out) //fills out[1..inp.length]  
  out(0) = a0 //prepends a0  
}
```



# Data parallelism

- So far, we have considered **task-parallelism**:
  - The `parallel(e1,e2)` primitive **distributes tasks** (i.e. the evaluation of its parameters) on parallel computing nodes
- We now move to **data-parallelism**:
  - Data are distributed on parallel computing nodes, and **on each datum** a **(parameterized) parallel task** is executed

```
def initializeArrayPar(xs: Array[Int])(v: Int): Unit = {  
  for (i <- (0 until xs.length).par) {  
    xs(i) = v  
  }  
}
```

- The data in the parallel collection are “split” in independent parts
- Parallel tasks act on the splitted parts of the collection
- A dedicated scheduler manages this parallelism

# Side effects

- The parallel tasks are not “functional”
  - they affect the program through **side-effects**
- As long as the data-parallel tasks write to **separate memory locations**, the program is correct
  - parallel write to the same location(s) **needs synchronization** to avoid data corruption

```
val mutual = new AnyRef()
def parMultiIncrement(n: Int) = {
  var v: Long = 0
  for (i <- (0 until n).par)
    mutual.synchronized{v=v+1}
  v
}
```

# Example

- Depict the **Mandelbrot** set

```
def run(n:Int, level:Int) : Unit = {  
  val out = new FileOutputStream(fileName)  
  out.write(("P5\n"+n+" "+n+"\n255\n").getBytes())  
  for (j <- (0 until n*n))  
  { val x = -2.0 + (j%n)*3.0/n  
    val y = -1.5 + (j/n)*3.0/n  
    var z = new Complex(0,0)  
    var c = new Complex(x,y)  
    var i = 0  
    while (z.abs < 2 && i < level)  
      {z = z*z + c; i=i+1}  
    out.write(255*(level-i)/level) }  
  out.close()  
}
```

# Example

- Depict the **Mandelbrot** set (parallel version)

```
def runPar(n:Int, level:Int) : Unit = {  
  val out = new FileOutputStream(fileName)  
  out.write(("P5\n"+n+" "+n+"\n255\n").getBytes())  
  var a=new Array[Int](n * n)  
  for (j <- (0 until n*n).par)  
  { val x = -2.0 + (j%n)*3.0/n  
    val y = -1.5 + (j/n)*3.0/n  
    var z = new Complex(0,0)  
    var c = new Complex(x,y)  
    var i = 0  
    while (z.abs < 2 && i < level)  
      {z = z*z + c; i=i+1}  
    a(j) = 255*(level-i)/level }  
  for{k <- 0 until n*n} out.write(a(k))  
  out.close()  
}
```

# Mandelbrot: evaluation

- Data-parallelism is **appropriate** in this case because:
  - Data-parallel tasks have different **workloads**
    - given the parallel task for datum  $i$ , let  $w(i)$  be its workload, i.e. the amount of work required to process it
  - For Mandelbrot we have  $w(i) = \text{\#iterations}$ 
    - Hence  $w(i)$  is not known a-priori, hence the programmer cannot appropriately manage the parallel executions (with primitives like **parallel**)
    - In fact, the complexity of parallel is the max of the two activities, and in case of **unbalanced** workload it is not efficient
  - Goal of the data-parallel scheduler:
    - Dynamically **balance** the workload across processors without any a-priori knowledge about the  $w(i)$

# Parallelizable collections

- Sequential collections can be parallelized using `.par`
  - Notice that some collections do not have a parallel counterpart

```
val vector = Vector.fill(10000000)('')  
val list = vector.toList  
vector.par // creates a ParVector[String]  
list.par   // also creates a ParVector[String]
```

- The reason is that their representation does not naturally support data split
- Array, Range and Vector are parallelizable (`ParArray[T]`, `ParRange[T]` and `ParVector[T]`)
- Other collections, like `Lists`, are not directly parallelizable; `.par` returns the closest parallelizable supertype in the hierarchy

# Exercise

- Compute set intersection:

```
def intersection(a: Set[Int], b: Set[Int]):  
  Set[Int] = {  
    val result = mutable.Set[Int]()  
    for (x <- a) if (b contains x) result += x  
    result  
  }  
intersection((0 until 10000).toSet,  
            (0 until 10000 by 4).toSet)
```

- Can we parallelize this function?

# Exercise

- Compute set intersection:

```
def intersection(a: Set[Int], b: Set[Int]):  
  Set[Int] = {  
    val result = mutable.Set[Int]()  
    for (x <- a.par) if (b contains x) result += x  
    result  
  }  
intersection((0 until 10000).toSet,  
            (0 until 10000 by 4).toSet)
```

- Can we parallelize this function?
  - Yes, BUT the parallel version is **not correct!**
  - PROBLEM: **concurrent** modification on the same (non thread-safe) data structure (result)



# Exercise

- Compute set intersection:

```
def intersection(a: Set[Int], b: Set[Int]):  
  ParSet[Int] = {  
    if (a.size < b.size) a.par filter (b contains _)  
    else b.par filter (a contains _)  
  }  
intersection((0 until 10000).toSet,  
            (0 until 10000 by 4).toSet)
```

- Here is a **correct** parallel version:
  - Exploits a (parallel) method directly provided by the parallel collection

# Concurrent read/writes

- Consider the following example:

```
val graph = mutable.Map[Int, Int]() ++  
  (0 until 100000).map(i => (i, i + 1))  
graph += (graph.size - 1) -> 0  
for ((k, v) <- graph.par) graph += k -> graph(v)  
graph.count( p =>  
  (p._2 != ((p._1 + 2) % graph.size)) )
```

- It contains two errors!
  - Concurrent tasks **modify** the non thread-safe Map collection
  - Concurrent tasks **read** from the collection while they modify it

# Concurrent read/writes

- Consider the following example:

```
val graph =  
    concurrent.TrieMap[Int, Int]() ++  
        (0 until 100000).map(i => (i, i + 1))  
graph += (graph.size - 1) -> 0  
val previous = graph.snapshot()  
for ((k, v) <- graph.par) graph += k -> previous(v)  
graph.count(p => (p._2 != ((p._1 + 2) % graph.size)))
```

- Here is a correct version:
  - TrieMap is a **thread-safe** collection (it supports concurrent modifications)
  - The method **snapshot** efficiently grab the current state