

Distributed Data-Parallel Programming

Distributed Data-Parallel Programming

- We have studied **parallel collections**
 - parallel tasks running in the same machine **sharing the same memory**
 - we have seen the impact on programming of **concurrent access to memory**
 - Bottleneck, data corruption, ...
- We now move to a **distributed scenario**
 - Data can be partitioned and distributed on **different machines** (that do not share memory)
 - Two main new concerns:
 - Partial **failures**: one node crashes, but the entire computation should not
 - Network **latency**: inter-node interaction is orders of magnitude slower

Google MapReduce

- In early 2000, Google proposed a novel programming model (supported by a corresponding framework) for **data processing on large clusters**:

MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

Google, Inc.

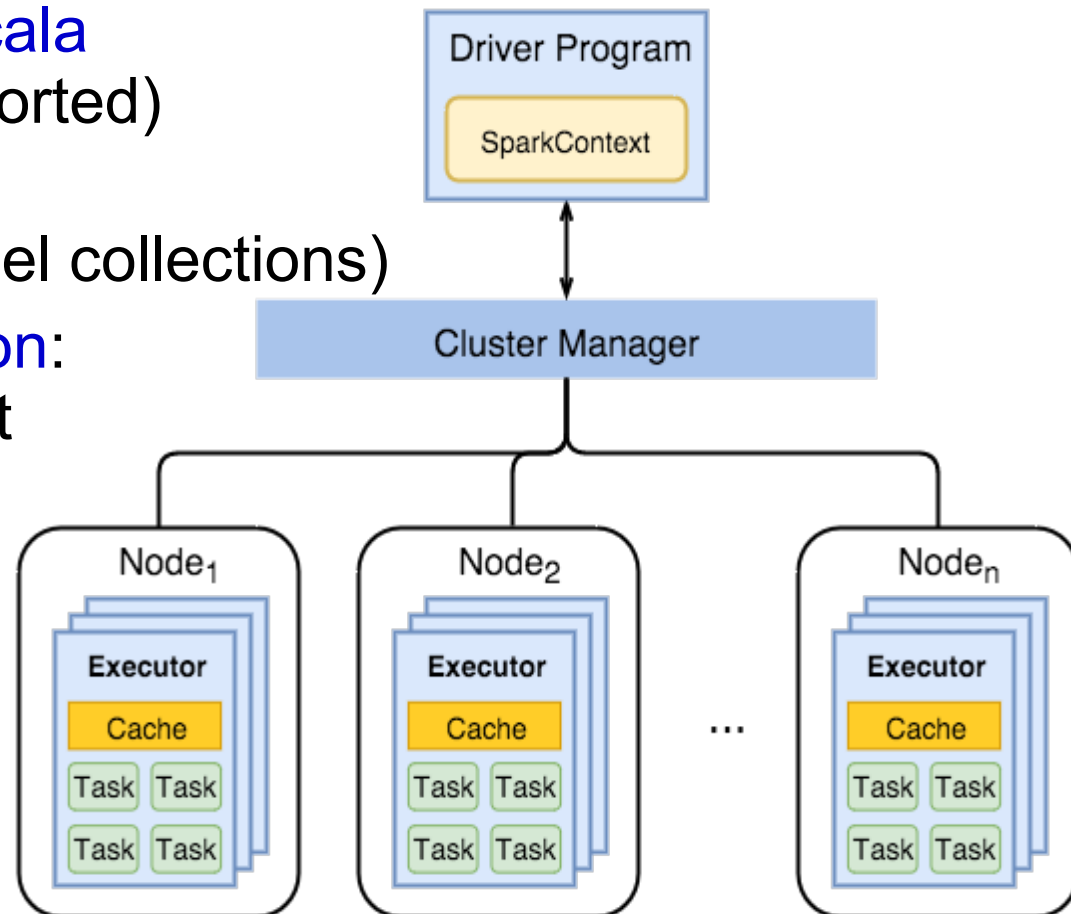
- Main properties:
 - **The framework exposes an easy to use interface**
(programs are expressed as map-reduce steps)
 - **Fault tolerant**
- **Apache Hadoop** is an open-source implementation of a map-reduce framework

Pros-Cons of Hadoop

- **Fault tolerance** permitted to run big-data applications on thousands of nodes
 - probability of one-node failure is **not negligible**
- But, fault tolerance comes at a price:
 - Between each map-reduce, data are always **stored** on disk
 - In case of failure, they can be **restored**
- Disk storage have a negative impact on **performances**
 - **Apache Spark** significantly improves Hadoop:
 - **Immutable** datasets
 - Intermediary data **stored in memory**
 - **Modifications** to the datasets **are recomputed** in case of failure

Spark

- **Spark** makes distributed data parallel programming more efficient and easy:
 - Natural to be used within **Scala** (Python and Java also supported)
 - **Higher-order** programming (as for collections and parallel collections)
 - It provides the **RDD collection**: Resilient Distributed Dataset
 - Data in an RDD are **distributed** over nodes:
 - The programmer refers to the entire distributed dataset as it was a **unique collection**



How to create an RDD

- New RDDs are created by exploiting a `SparkContext` object:
 - It is the handle to the Spark cluster (connection between the cluster and the running application)
 - Two main methods:
 - `parallelize`: convert a local collection to an RDD
 - `textFile`: read a text file, return an `RDD[String]` (one String per line)
- Another way to create an RDD is:
 - Apply higher-order functions called `transformations` (similar to map or filter) to an existing RDDs

```
val largelist: List[String] = ...  
val wordsRdd = sc.parallelize(largelist)  
val lengthsRdd = wordsRdd.map(_.length)
```

A first example

- **WordCount** is considered the HelloWorld for distributed data parallel programming
 - This is the Scala+Spark version:

```
val input = sc.textFile(inputFile)
val counts = input.flatMap(line => line.split(" ")).
  map(w => (w.filter(_._isLetter).toUpperCase, 1)).
  reduceByKey((x, y) => x + y)
```

- Following the programming model initiated by Google MapReduce, it is based on a **map-reduce** computation
 - The data elaboration is executed in parallel on each data partition, and **only the local reduce results are communicated** and combined

Higher-order functions on RDD

- The functions available for computations on RDDs are classified in:
 - **Transformations**: return a new RDD
 - Examples are `map`, `flatMap`, `filter` and `distinct`
 - **Actions**: do not return an RDD
 - Examples are `collect`, `count`, `take`, `reduce`, `foreach`, `takeSample`, `takeOrdered` and `saveAsTextFile`
 - There are also transformations on two RDDs:
 - Examples are `union`, `intersection`, `subtract` and `cartesian`

IMPORTANT NOTE:

Transformations are lazy, i.e., evaluated only when necessary to compute a subsequent action

Transformations signature

`map[B](f: A=> B): RDD[B]`

`flatMap[B](f: A=> Traversable[B]): RDD[B]`

`filter(pred: A=> Boolean): RDD[A]`

`distinct(): RDD[B]` (remove duplicates)

- Transformations on two RDDs

`union(other: RDD[T]): RDD[T]`

(do not remove duplicates)

`intersection(other: RDD[T]): RDD[T]`

(remove duplicates)

`subtract(other: RDD[T]): RDD[T]`

`cartesian[U](other: RDD[U]): RDD[(T, U)]`

Actions signature

`collect(): Array[T]`

`count(): Long`

`take(num: Int): Array[T]`

`reduce(op: (A, A) => A): A`

`fold(z: A)(f: (A, A) => A): A`

`aggregate[B](z: B)(seqop: (B, A) => B,
 combop: (B, B) => B): B`

`foreach(f: T => Unit): Unit`

`takeSample(withRepl: Boolean, num: Int): Array[T]`

`takeOrdered(num: Int)(
 implicit ord: Ordering[T]): Array[T]`

`saveAsTextFile(path: String): Unit`

Lazy evaluation of transformations

- Laziness of transformations is at the basis of the improved Spark **performances**:
 - Spark **do not compute immediately** the data resulting from transformations
 - Deferring transformations allows Spark to **analyse** and **optimize** the chain of operations before executing it:
 - All transformations are stored in a **DAG** (direct acyclic graph):
 - Transformations are organized in **stages**
 - Each stage contains a sequence of “**narrow**” transformations (like **map** or **filter**)
 - “**wide**” transformations (like **intersection**) start new stages
 - Narrow transformations can be executed **locally** in each partition
 - Wide transformations needs **shuffling** (inter-partition data exchange); these represent stage boundaries
- Evaluation of transformations are triggered by **actions**

Pros of Laziness

- Deferring the evaluation of transformations allows Spark:
 - to compute **only what is needed**
 - If an action takes only few data, only those strictly needed are computed

```
val lastYearslogs: RDD[String] = ...  
val firstlogsWithErrors =  
lastYearslogs.filter(_.contains("ERROR")).take(10)
```

- to minimize data **traversal**
 - Narrow transformations can be pipelined to be computed in **one pass**

```
time(words.filter(w => w.length>10).count())  
val pairs = words.map(w => w.length)  
val longWords = pairs.filter(l => (l > 10))  
time(longWords.count())
```

Cons of Laziness

- Transformations could be computed several times:
 - If more than one action depends on the same transformation, it is **re-computed** for each action
 - This problem can be avoided by asking to make the intermediary data **persistent**

```
val counts = words.  
  map(word =>  
    (word.filter(_._isLetter).toUpperCase, 1)).  
  reduceByKey((x, y) => x + y).persist()  
counts.saveAsTextFile(outputFile)  
counts.reduce(( p1: (String,Int),  
               p2:(String,Int)) =>  
  if (p1._2 > p2._2) p1 else p2))
```

Persistence

- Persistence can be customized by passing to the `persist()` method the **storage level** (add `_n` for n replications):

<i>Level</i>	<i>Space used</i>	<i>CPU time</i>	<i>In memory</i>	<i>On disk</i>
MEMORY_ONLY	High	Low	Y	N
MEMORY_ONLY_SER	Low	High	Y	N
MEMORY_AND_DISK*	High	Medium	Some	Some
MEMORY_AND_DISK_SER [†]	Low	High	Some	Some
DISK_ONLY	Low	High	N	Y

```
import org.apache.spark.storage.StorageLevel
val counts = words.
  map(word =>
    (word.filter(_.isLetter).toUpperCase, 1)).
  reduceByKey((x, y) => x + y).
  persist(StorageLevel.DISK_ONLY_2)
```

Application vs Executor Evaluation

- Transformations are **evaluated by executors** inside the computing nodes
 - Narrow locally, wide with shuffling
- Actions **return results to application** through the driver

```
case class Person (name: String age: Int)
val people: RDD[Person] = ...
people.foreach(println) //executed by the executor
people.collect().foreach(println)//by the application
```

Key-Value Pairs

- The programming model initiated by MapReduce used **key-value** pairs as basic way to store intermediary data:

As a reaction to this complexity, we designed a new abstraction that allows us to express the simple computations we were trying to perform but hides the messy details of parallelization, fault-tolerance, data distribution and load balancing in a library. Our abstraction is inspired by the *map* and *reduce* primitives present in Lisp and many other functional languages. We realized that most of our computations involved applying a *map* operation to each logical “record” in our input in order to compute a set of intermediate key/value pairs, and then applying a *reduce* operation to all the values that shared the same key, in order to combine the derived data appropriately. Our use of a functional model with user-specified map and reduce operations allows us to parallelize large computations easily and to use re-execution as the primary mechanism for fault tolerance.

Key-Value RDDs

- Key-Value pairs are present also in Spark:
 - An `RDD[(K,V)]` is a collection of key (of type K) - value (of type V) pairs
 - Useful for `parallel` computation on distinct keys and for `regrouping` data (i.e. all pairs with the same key located in the same partition)
 - Additional `transformations/actions` for key-value RDDs:
 - `groupByKey(): RDD[(K, Iterable[V])]`
 - `reduceByKey(func: (V,V)=> V): RDD[(K, V)]`
 - `join[W]: (other: RDD[(K, W)]): RDD[(K, (V, W))]`
 - `leftOuterJoin[W](other: RDD[(K, W)]): RDD[(K, (V, Option[W]))]`
 - `rightOuterJoin[W](other: RDD[(K, W)]): RDD[(K, (Option[V], W))]`

Shuffling

- Operations like `groupByKey` group all pairs with the same key in a unique new pair
 - Moving pairs to the same node is necessary: this movement is called `shuffling`
 - Shuffling is the `most expensive` activity
- The programmer should `minimize shuffling` being aware of the shuffling performed by the functions:
 - E.g. `reduceByKey` computes local intermediary data and communicates across the network only such data

Example

- Compute the **number** and the **total amount** of the purchases of each customer:

```
case class Purchase(customerId: Int, price: Double)
val purchasesRdd: org.apache.spark.rdd.RDD[Purchase] =
    sc.textFile(inputFile)

val purchasesPerCustomer: Array[(Int, (Int, Double))] =
    purchasesRdd.
    map(p => (p.customerId, p.price)).    // Pair RDD
    groupByKey().    // returns RDD[(K, Iterable[V])]
    map(p => (p._1, (p._2.size, p._2.sum))).
    collect()
```

Example

- Compute the **number** and the **total amount** of the purchases of each customer:

```
case class Purchase(customerId: Int, price: Double)
val purchasesRdd: org.apache.spark.rdd.RDD[Purchase] =
    sc.textFile(inputFile)

val purchasesPerCustomer: Array[(Int, (Int, Double))] =
    purchasesRdd.
    map(p => (p.customerId, p.price)).    // Pair RDD
    groupByKey().    // returns RDD[(K, Iterable[V])]
    map(p => (p._1, (p._2.size, p._2.sum))).
    collect()
```

- This version is **not efficient** because it shuffles all ill-placed pairs

Example

- Compute the **number** and the **total amount** of the purchases of each customer:

```
case class Purchase(customerId: Int, price: Double)
val purchasesRdd: org.apache.spark.rdd.RDD[Purchase] =
    sc.textFile(inputFile)

val purchasesPerCustomer: Array[(Int, (Int, Double))] =
    purchasesRdd.
    map(p => (p.customerId, (1, p.price))). //Pair RDD
    reduceByKey((v1, v2) =>
        (v1._1 + v2._1, v1._2 + v2._2)).
    collect()
```

- **More efficient:** shuffles at most one pair per key from each node

Partitioning

- In pair RDDs, it is possible to **customize** distribution of pairs over the partitions as a function on the keys
 - Two available library **functions**:
 - **RangePartitioner**: partitions computed dividing the key domain (from minimum to maximum) in a given number of equally sized ranges
 - **HashPartitioner**: applies an hash function to keys and distributes according to the hash modulo the number of partitions

```
val purchasesPairs = purchasesRdd.  
  map(p => (p.customerId, (1, p.price)))  
val purchasesRangeP = purchasesPairs.  
  partitionBy(new RangePartitioner(numPartitions,  
                                     purchasesPairs))  
  
val purchasesHashP = purchasesPairs.  
  partitionBy(new HashPartitioner(numPartitions))
```

Example

- Using a partitioner in the customer example, **no shuffling** occurs:

```
case class Purchase(customerId: Int, price: Double)
val purchasesRdd: org.apache.spark.rdd.RDD[Purchase] =
    sc.textFile(inputFile)

val purchasesPairs = purchasesRdd.
    map(p => (p.customerId, (1, p.price)))
val purchasesPartitioned = purchasesPairs.
    partitionBy(new RangePartitioner(numPartitions,
                                     purchasesPairs))
val purchases = purchasesPartitioned.
    reduceByKey((v1, v2) =>
        (v1._1 + v2._1, v1._2 + v2._2)).
    collect()
```

Implicit use of partitioner

- There are functions that hold (or propagate) a partitioner:

- ▶ `cogroup`
- ▶ `groupWith`
- ▶ `join`
- ▶ `leftOuterJoin`
- ▶ `rightOuterJoin`
- ▶ `groupByKey`
- ▶ `reduceByKey`
- ▶ `foldByKey`
- ▶ `combineByKey`
- ▶ `partitionBy`
- ▶ `sort`
- ▶ `mapValues` (if parent has a partitioner)
- ▶ `flatMapValues` (if parent has a partitioner)
- ▶ `filter` (if parent has a partitioner)

- E.g. `mapValues` is preferable to `map` in case of partitioned RDDs because the partitioner is preserved

Example: PageRank

- Famous algorithm (named after Larry Page, co-founder of Google) used to rank importance of websites
- Consider a graph $G=(V,E)$, with arcs $(v,k) \in E$ meaning that document v has a link to document k
- Iterative algorithm that associates to each document k a rank $rank(k)$
 - In the first iteration $rank_0(k) = 1.0$, for all k
 - At each iteration, $rank_{i+1}(k) = 0.15 + 0.85 * c_i(k)$, where
$$c_i(k) = \sum_{(v,k) \in E} rank_i(v) / |\{t \mid (v,t) \in E\}|$$

PageRank implementation

```
val links = edges.partitionBy(new HashPartitioner(4)).
    groupByKey().persist()
var ranks = links.mapValues(v => 1.0)

for(i <- 0 until 10) {
    val contributions = links.join(ranks).flatMap {
        case (u, (uLinks, urank)) =>
            uLinks.map(t => (t, urank / uLinks.size))
    }
    ranks = contributions.
        reduceByKey((x,y) => x+y).
        mapValues(v => 0.15+0.85*v)
}

ranks.saveAsTextFile(outputFile)
```

Scalability

- How much faster can a given problem be solved with n nodes instead of one?
- How much more work can be done with n nodes instead of one?
- What impact for the communication requirements of the distributed application have on performance?
- What fraction of the resources is actually used productively for solving the problem?

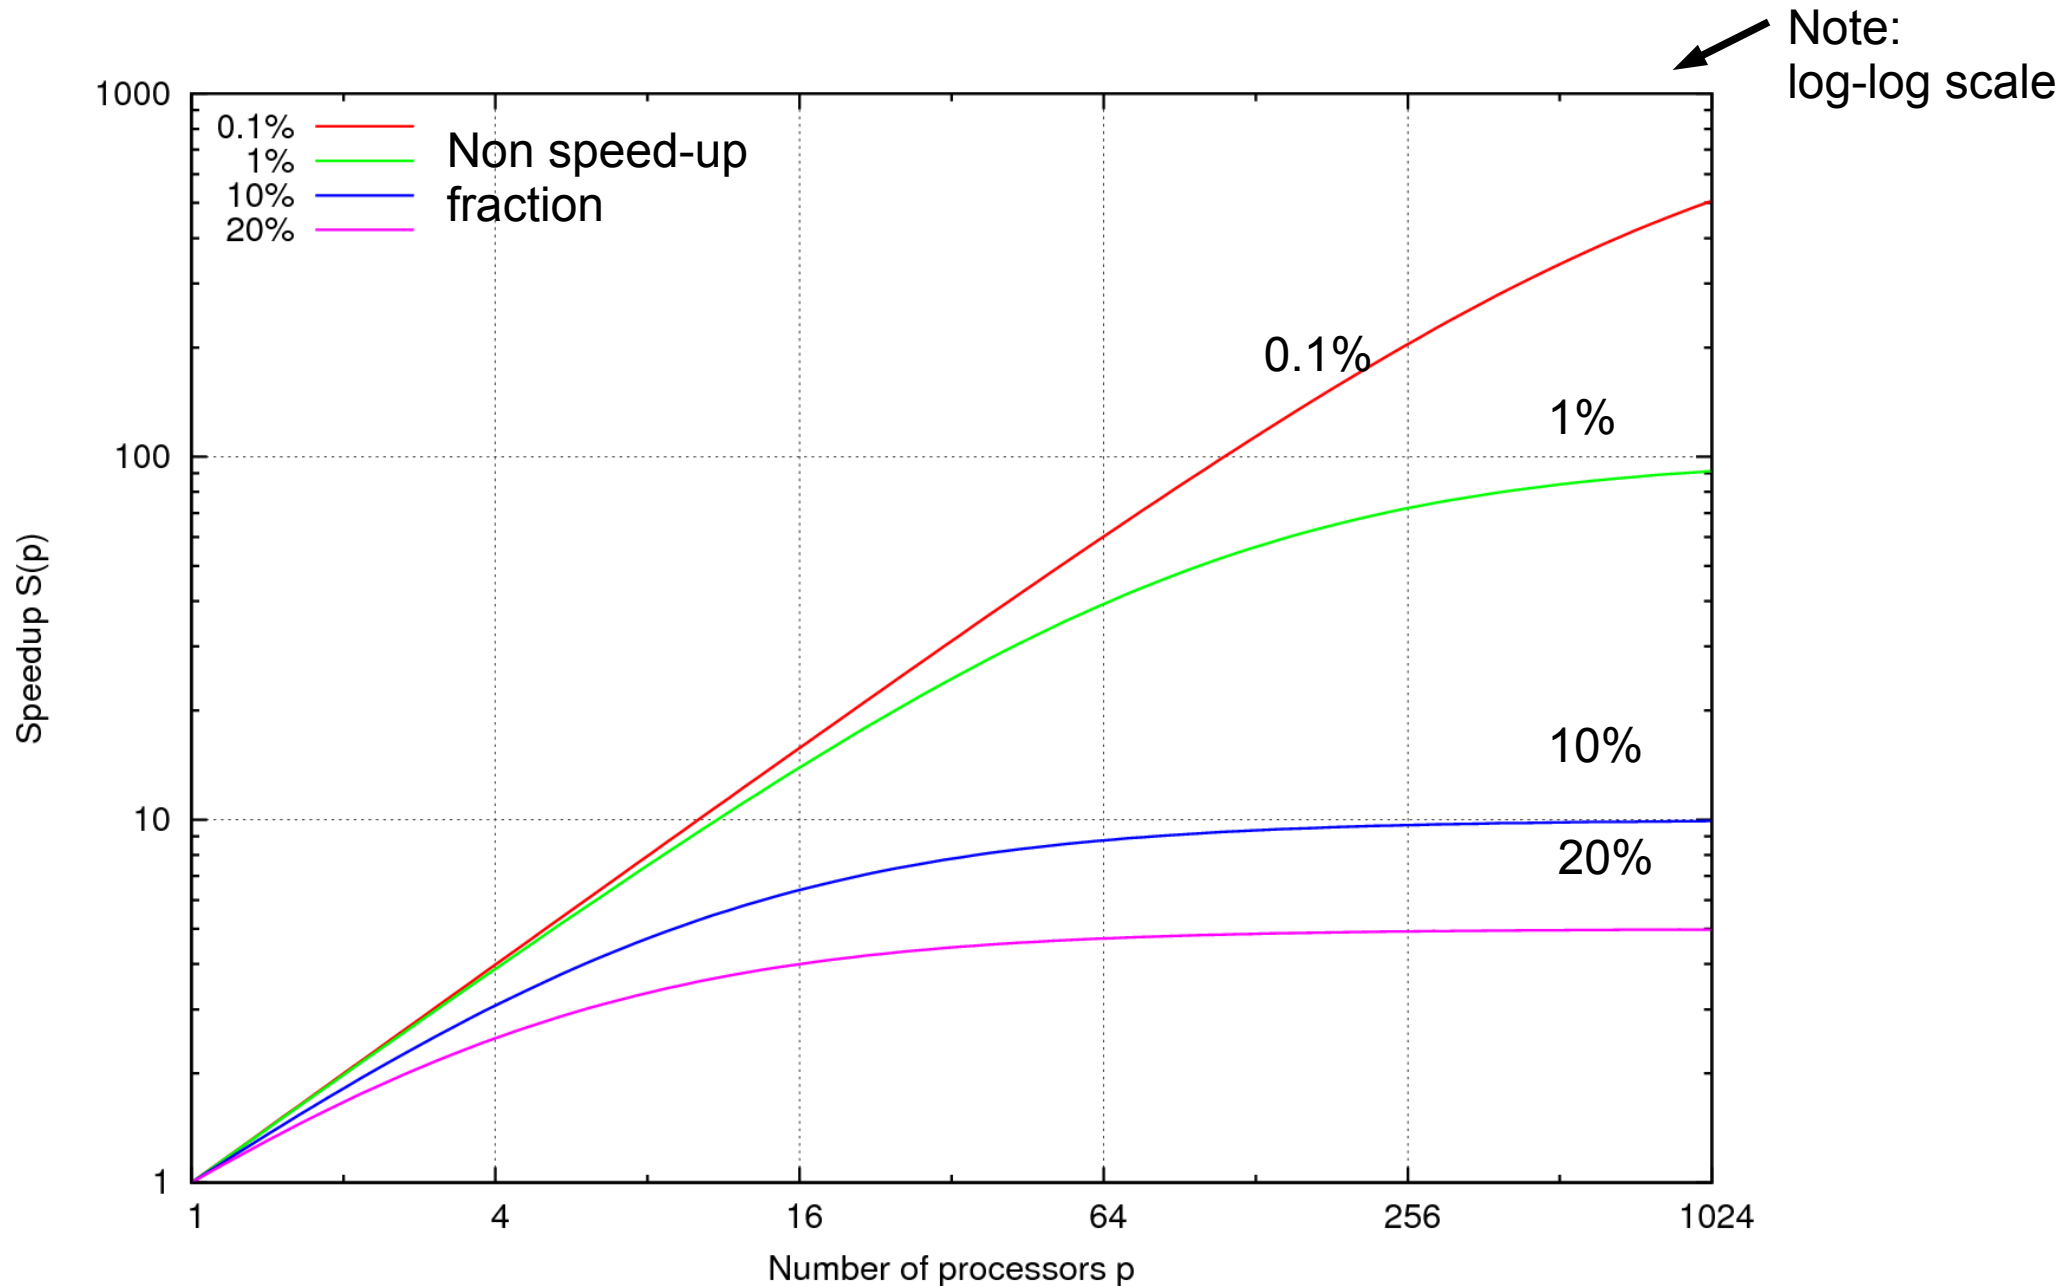
Speedup

- Let us define:
 - n = Number of nodes
 - $T(n)$ = Execution time of the parallel program with n nodes
- Speedup $S(n)$

$$S(n) = \frac{T(1)}{T(n)}$$

- In the ideal case, the program with n nodes requires $1/n$ the time of the program with 1 node
- $S(n) = n$ is the ideal case of linear speedup
 - Realistically, $S(n) \leq n$

Speedup (according to Amdahl's law)



Scaling Efficiency

- *Objective:*
 - Evaluate the impact of Amdahl's law on your distributed program
 - Quantify the effect on the execution time for each node that is added
- *Solution:* measure *Strong Scaling Efficiency*
 - Increase the number of nodes *n* keeping the *total* problem size fixed
 - The total amount of work remains constant, while the amount of work for each processor decreases as *n* increases
 - How to quantify the impact of each added node?
 - Divide the speedup for the number of nodes

Strong Scaling Efficiency


- $E(n)$ = Strong Scaling Efficiency

$$E(n) = \frac{S(n)}{n} = \frac{T(1)}{n T(n)}$$

where

- $T(n)$ = Execution time of the execution with n nodes

Strong Scaling Efficiency and Amdahl's Law



Note: logarithmic x scale

Negative result on strong scaling efficiency

- The **strong scaling efficiency** always tends to zero because the **speedup** is limited by the constant $1/f$
- But in many cases we want to add computing resources to **do more computational work** instead of reducing the execution time for a fixed input
 - The goal is to solve larger problems within the same amount of time by increasing the computational power accordingly
 - How can we check whether adding computational power is actually effective to achieve this goal?

Weak Scaling Efficiency

- An alternative measure that considers increasing problem size is *Weak Scaling*:
 - Increase the number of nodes n keeping the per-node work fixed
 - The total amount of work grows as n increases
- $W(n)$ = Weak Scaling Efficiency

$$W(n) = \frac{T_1}{T_n}$$

where

- T_1 = time required to complete 1 work unit with 1 node
- T_n = time required to complete n work units with n nodes
- Weak scaling efficiency = 1 means that our program uses productively 100% of the added resources

Weak Scaling Efficiency

- **Question:** given the size of the input for 1 node (1 work unit) how much the input should be increased for n nodes (n work units)?
 - Let m be the input for 1 node: 1 work unit = $W(m)$ where $W(m)$ is the work complexity
 - Consider now n nodes: n work units = $n W(m)$ we need to find k such that $W(k m) = n W(m)$
- **Examples:**
 - Linear work complexity: $W(m) = m$
 $W(k m) = n W(m)$ implies $k m = n m$ hence $k = n$
 - Quadratic work complexity: $W(m) = m^2$
 $W(k m) = n W(m)$ implies $k^2 m^2 = n m^2$ hence $k = n^{1/2}$
 - Cubic work complexity: $W(m) = m^3$
 $W(k m) = n W(m)$ implies $k^3 m^3 = n m^3$ hence $k = n^{1/3}$